

CHAPTER 9 HEAP STRUCTURES

9.1 Min-max heaps

- Definition

- A *double-ended priority queue* is a data structure that supports the following operations:
 - (1) **Insert** an element with arbitrary key
 - (2) **Delete** an element with the **largest** key
 - (3) **Delete** an element with the **smallest** key
- When only insertion and one of the two deletion operations is supported, we may use a **min heap** or **max heap**. A **min-max heap** however supports all of the operations just described.

– Definition:

A *mix-max heap* is a **complete binary tree** such that if it is not empty, each element has a field called *key*. Alternating levels of this tree are min levels and max levels, respectively. Let x be any node in a min-max heap. **If x is on a min level** then the element in x has the minimum key from among all elements in the subtree with root x . We call this node a *min* node. Similarly, **if x is on a max level** then the element in x has the maximum key from among all elements in the subtree with root x . We call this node a *max* node.

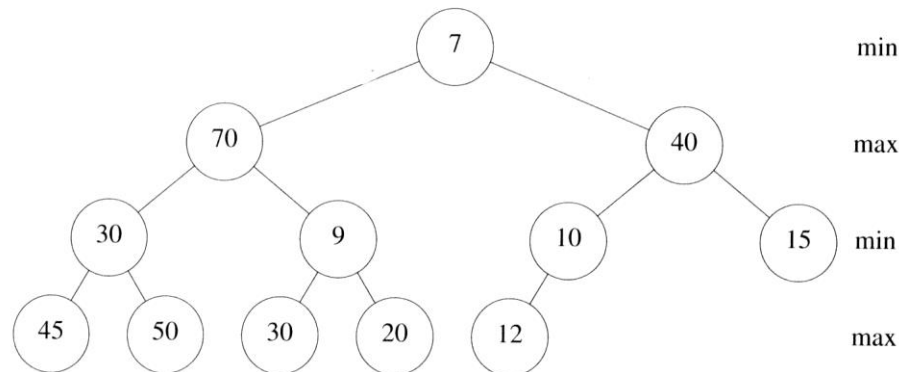


Figure 9.1: A 12 element min-max-heap

– Insertion into a min-max heap

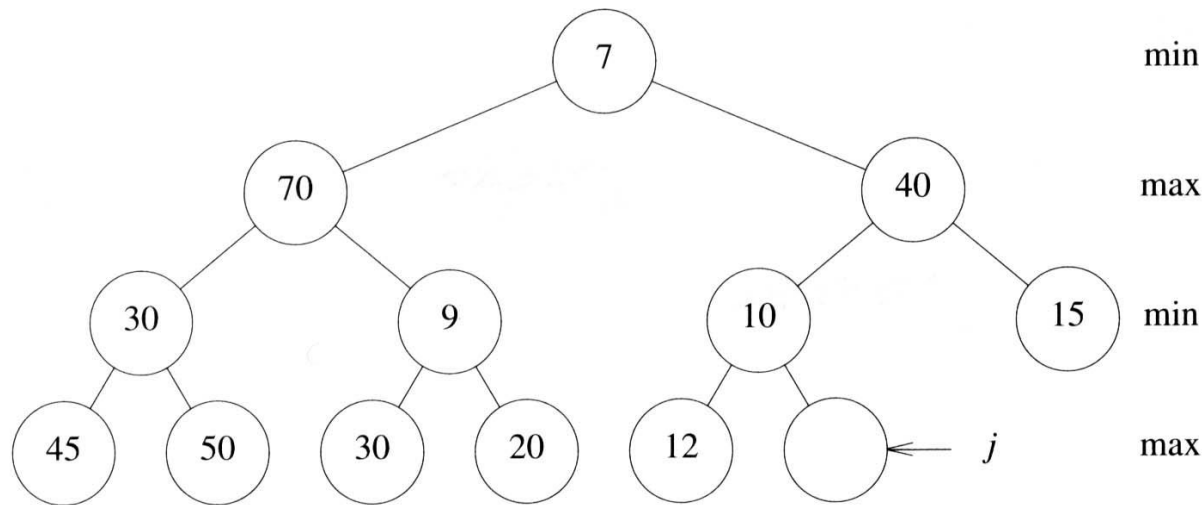
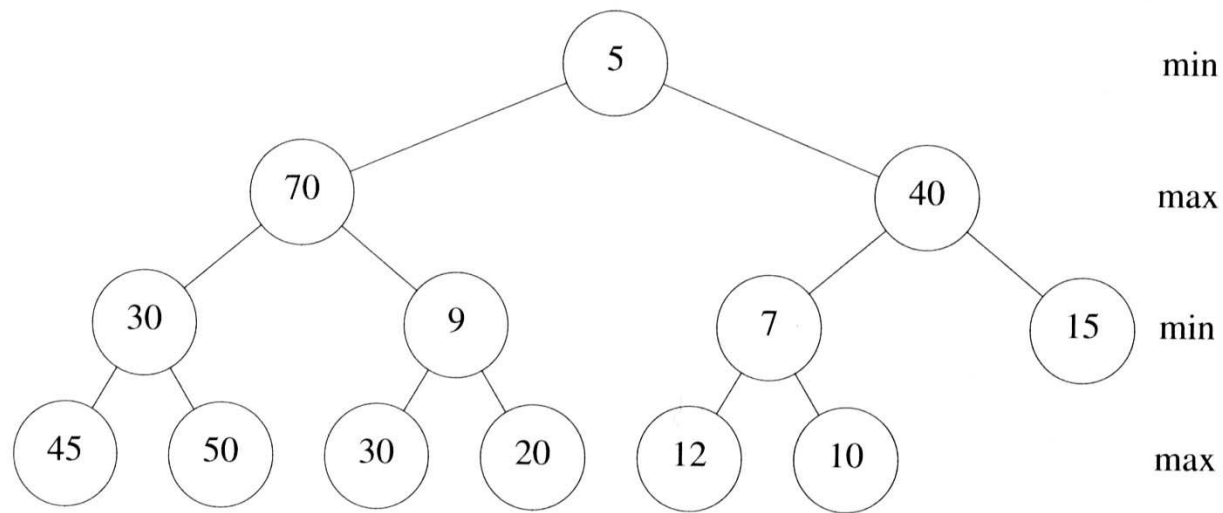


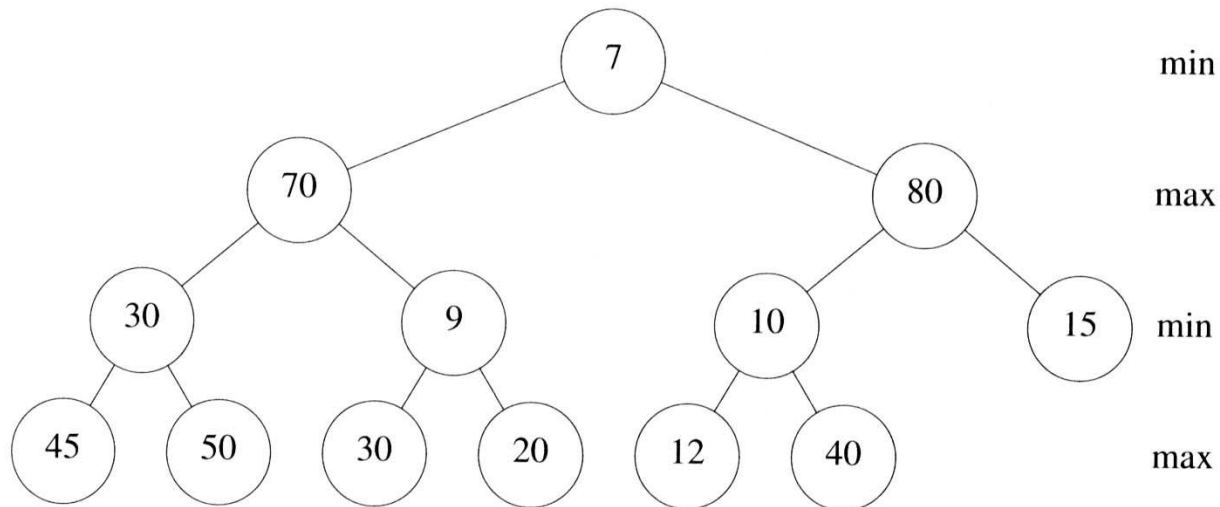
Figure 9.2: Min-max-heap of Figure 9.1 with new node j

Insertion at a “max” level: (Insertion at a “min” level 以類比作法)

1. If it is **smaller than its father** (a “min”), then it must be smaller than all “max” above. So simply check the “min” ancestors.
2. If it is **greater than its father** (a “min”), then it must be greater than all “min” above. So simply check the “max” ancestors.



(a) min-max-heap of Figure 9.1 after inserting 5



(b) min-max-heap of Figure 9.1 after inserting 80

- **Analysis of `min_max_insert`:**

- Since a min-max heap with n elements has $O(\log n)$ levels, the complexity of the *min_max_insert* function is $O(\log n)$.

– Deletion of min element

- If we wish to delete the element with smallest key, then this element is in the root.
- In the case of the min-max-heap of Figure 9.1, we are to delete the element with key 7.
 - Following the deletion, we will be left with a min-max-heap that has 11 elements. Its shape is shown in Figure 9.4.
 - The node with key 12 is deleted from the heap and the element with key 12 is reinserted into the heap.
 - As in the case of deletion from a min or max-heap, the reinsertion is done by examining the nodes of Figure 9.4 from the root down towards the levels.
- In general situation, we are to reinsert an element *item* into a min-max-heap, *heap*, whose root is empty. We consider the two cases:

- (1) *The root has no children.* In this case *item* is to be inserted into the root.
- (2) *The root has at least one child.* Now, the smallest key in the min-max-heap is in one of the children or grandchildren of the root. We determine which of these nodes has the smallest key. Let this be node *k*. The following possibilities need to be considered:
 - » *item.key* \leq *heap[k].key*. Item may be inserted into the root as there is no element in heap with key smaller than *item.key*.
 - » *item.key* $>$ *heap[k].key* and *k* is a child of the root. Since *k* is a max node, it has no descendants with key larger than *heap[k].key*. Hence, node *k* has no descendants with key larger than *item.key*. So, the element *heap[k]* may be moved to the root and *item* inserted into node *k*.
 - » *item.key* $>$ *heap[k].key* and *k* is a grandchild of the root. In this case too, *heap[k]* may be moved to the root. Let *parent* be the parent of *k*. If *item.key* $>$ *heap[parent].key*, then *heap[parent]* and *item* are to be interchanged. This ensures that the max node *parent* contains the largest key in the sub-heap with root *parent*. At this point, we are faced with the problem of inserting *item* into the sub-heap with root *k*. the root of this sub-min-max-heap

is presently empty. This is quite similar to our initial situation where we were to insert *item* into the min-max-heap *heap* with root 1 and node 1 is initially empty. Therefore, we repeat the above process.

Delete a “min” – i.e., root:

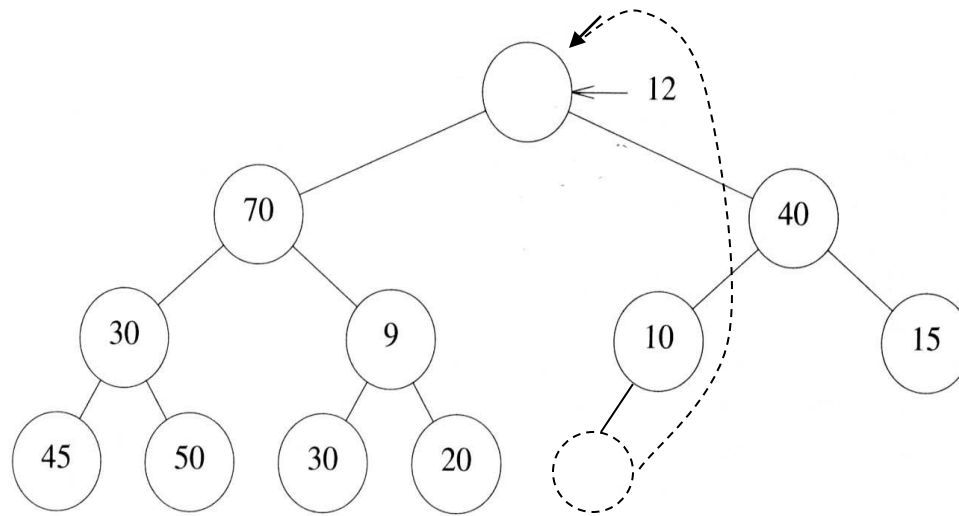


Figure 9.4: Shape of Figure 9.1 following a delete min

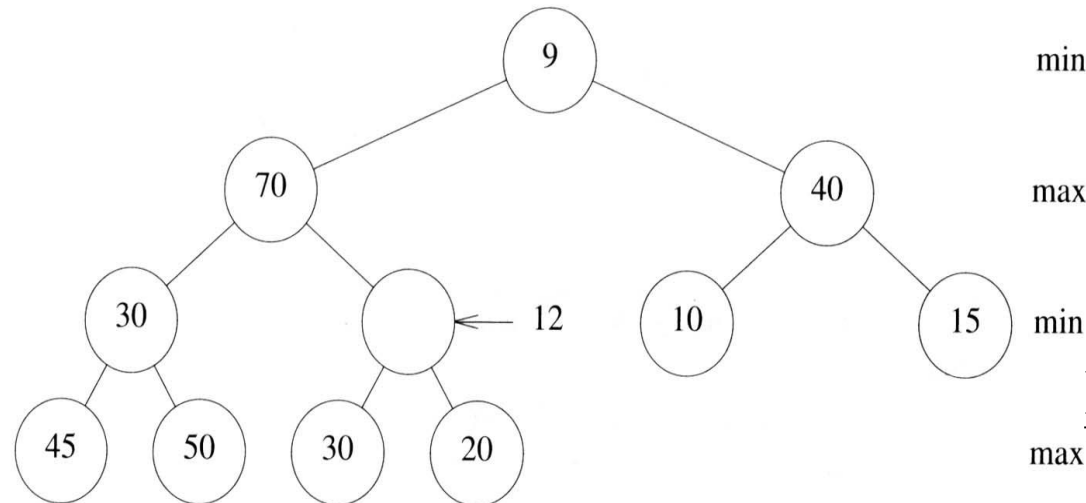


Figure 9.5: Figure 9.4 following the move of the element with key 9

min 1. Single node case.

2. 找出兒、孫層最小值 $\text{heap}[k].\text{key}$.

max (a) “12” $\leq \text{heap}[k].\text{key}$,
12直接放入 root

min (b) “12” $> \text{heap}[k].\text{key}$ and **k is a child**, 則 “12” 與 k 對調 (因為這case 一定只有一層(兒子), 沒有孫子層了。)

max (c) “12” $> \text{heap}[k].\text{key}$ and k is a **grandchild**, 則 “12” 與 k 對調, 之後, k不必去與其他的 min 再比對, 因 k 已經是兒孫層最小的了。

但 “12” 要與其下的 subtree 再做比較調整。調整的方式, 就是以 “12” 為root, 並與上面作一樣的動作即可。

Delete a “max”: 則由它以下之
max 兩個level找最大的, 然後與上面一樣作類似的調整。

- **Analysis of *delete_min*:**

- Since a min-max heap is a complete binary tree, *heap* has $O(\log n)$ levels. Hence, the complexity of *delete_min* is $O(\log n)$.

9.2 Deaps

- Definition
 - A **deap** is a **double-ended heap** that supports the double-ended priority queue operations of insert, delete min, and delete max.
 - As in the case of the min-max heap, these operations take logarithmic time on a deap. However, the **deap is faster** by a constant factor and **algorithms are simpler**.

– Definition:

- A *deap* is a **complete binary tree** that is either empty or satisfies the following properties:
 - The root contains no element.
 - The left subtree is a min-heap.
 - The right subtree is a max-heap.
 - If the right subtree is not empty, then **let i be any node in the left subtree. Let j be the corresponding node in the right subtree.** If such a j does not exist, then let j be the node in the right subtree that corresponds to the parent of i . The **key in node i is less than or equal to the key in j .**

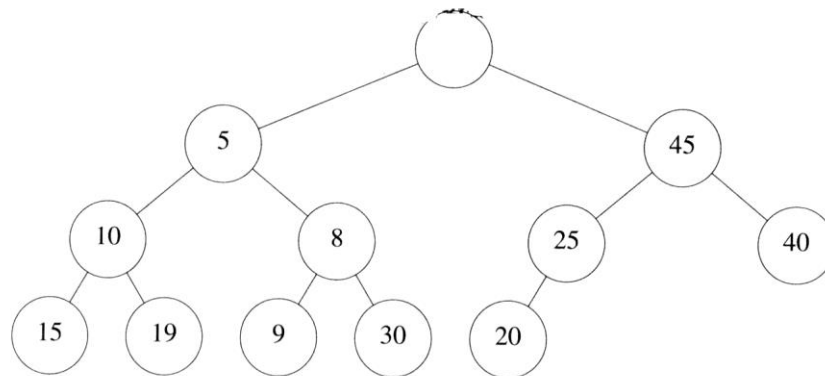


Figure 9.6: An 11 element deap

- In an n element deap, $n > 1$, the min element is in the root of the min-heap while the max element is in the root of max-heap.
- If $n = 1$, then the min and max elements are the same and are in the root of the min heap.
- Since a **deap is a complete binary tree**, it may be stored as an implicit data structure in a **one dimensional array**.
 - In the case of a deap, position 1 of the array is not utilized.
 - Let n denote the last occupied position in this array. Then the number of elements in the deap is $n-1$.
 - If **i is a node** in the min-heap, then its **corresponding node in the max-heap is $i+2^{\lfloor \log_2 i \rfloor - 1}$** . Hence the j defined in property (4) of the definition is given by:

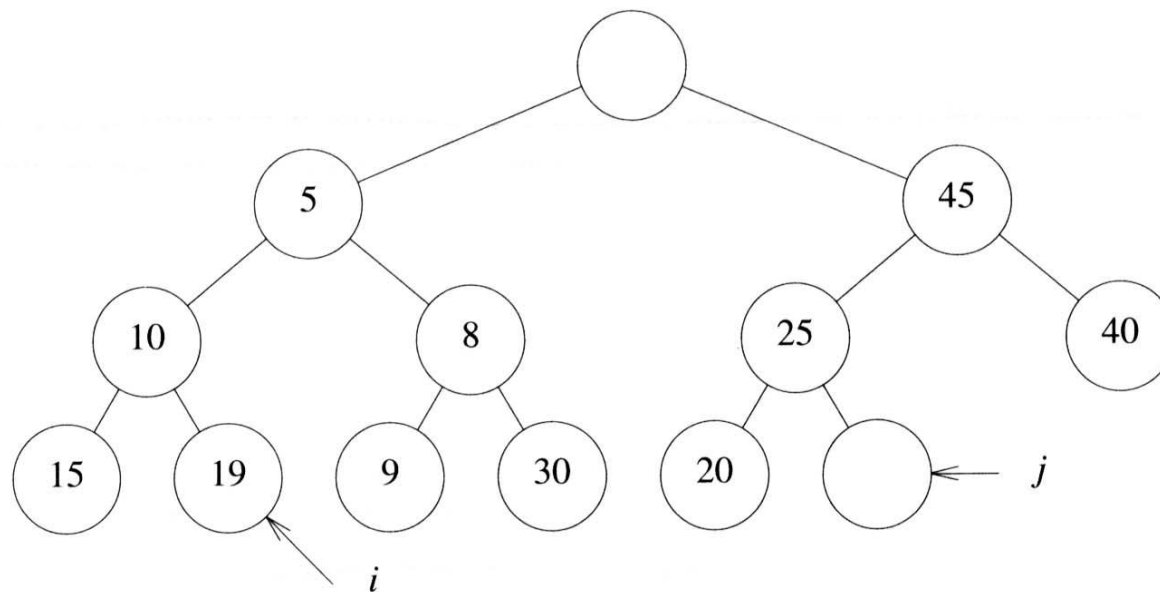
$$j = i + 2^{\lfloor \log_2 i \rfloor - 1};$$

$$\text{if } (j > n) \text{ } j /= 2;$$

- Notice that if property (4) of the deap definition is satisfied by all leaf nodes i of min-heap, then it is satisfied by all remaining nodes of the min-heap too.

- **Insertion** into a deap

- Suppose we wish to insert an element with key 4 into the deap of Figure 9.6. Following this insertion, the deap will have 12 elements in it and will thus have the shape shown in Figure 9.7. j points to the new node in the deap.



Insert 4 into
the deap

Figure 9.7: Shape of a 12 element deap

- The insertion process begins by comparing the key 4 to the key in j 's corresponding node, i , in the min-heap. The node contains a 19.
- To satisfy property (4), we move the 19 to node j . Now, if we use the min-heap insertion algorithm to insert 4 into position i , we get the deap of Figure 9.8.

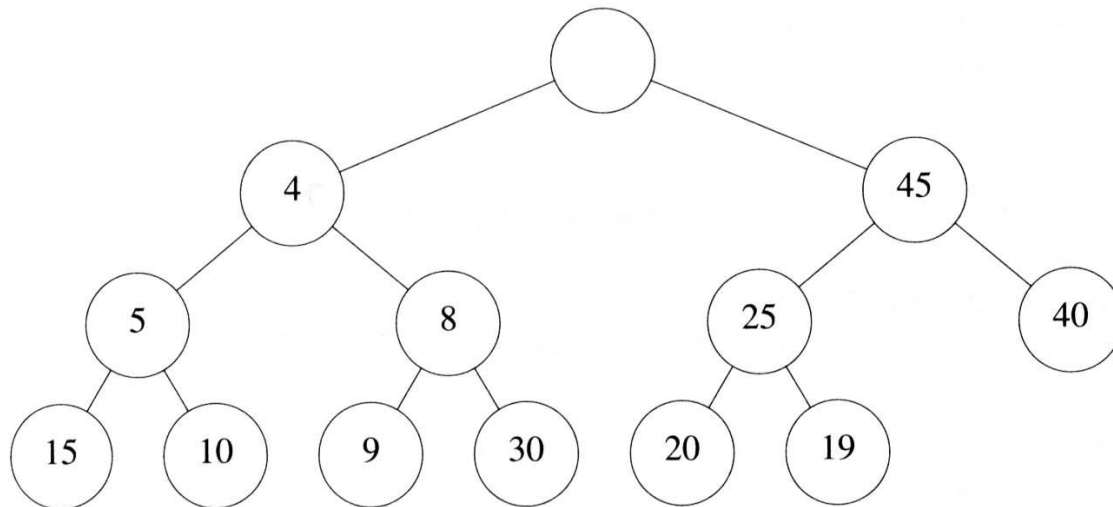
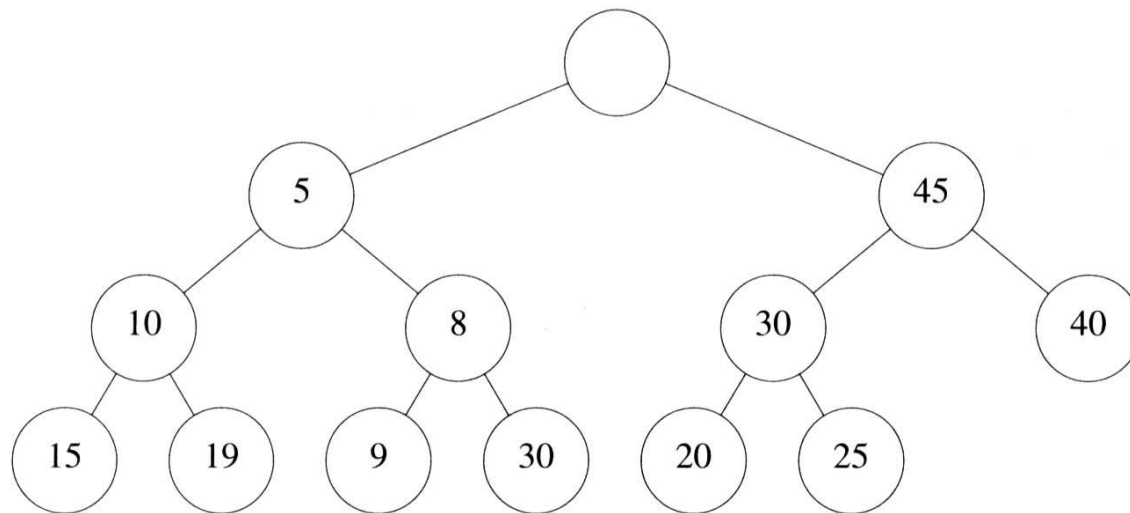


Figure 9.8: Deap of Figure 9.6 following the insertion of 4

- If instead of inserting a 4, we were to insert a 30 into the deap of Figure 9.6, then the resulting deap has the same shape as in Figure 9.7.
- Comparing 30 with the key 19 in the corresponding node i , we see that property (4) may be satisfied by using the max-heap insertion algorithm to insert 30 into position j . This results in the deap of Figure 9.9.



Insert 30 into
the deap of
Figure 9.7

Figure 9.9: Deap of Figure 9.6 following the insertion of 30

– **Analysis of *deap_insert*:**

- Time complexity of *deap_insert* is $O(\log n)$ as the the height of the deap is $O(\log n)$.

- **Deletion** of min element

- **Example:** Suppose that we wish to remove the minimum element from the deap of Figure 9.6.

- We **first place the last element** (the one with key 20 (Figure 9.7)) in the deap into a temporary element, temp, since the deletion removes this node from the heap structure.
- Next, we **fill the vacancy created in the min-heap root** (node 2) by the removal of the minimum element. To fill this vacancy we move along the path from the root to a leaf node.
 - » Prior to each move, we place **the smaller of the elements in the current node's children** into the current node. We then move to the node previously occupied by the moved element.
 - » In this example, we first move 8 to node 2. Then we move 9 into the node formerly occupied by 8.
- Now, We have an empty leaf and proceed to insert 20 into this. We **compare 20 with the key 40 in its max partner**. Since $20 < 40$, no exchange is needed and we proceed to insert 20 into the min-heap beginning at the empty position. This operation results in the deap of Figure 9.10.

» $\text{max_partner}(n)$. This function computes the max-heap node that corresponds to the parent of the min-heap position n .

$$\text{max_partner}(n) = (n + 2^{\lfloor \log_2 n \rfloor - 1}) / 2$$

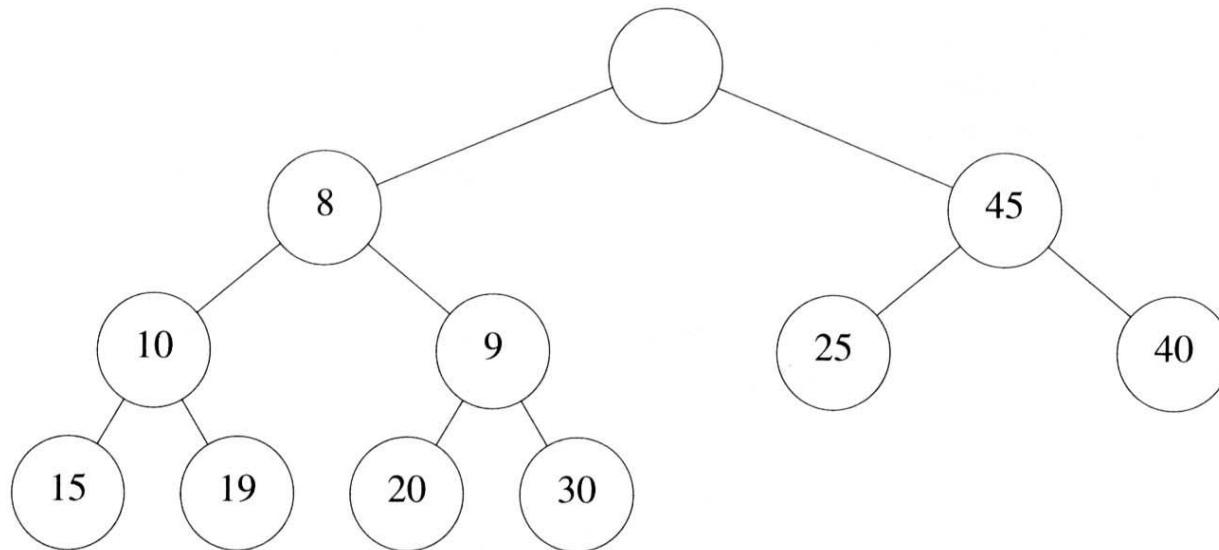


Figure 9.10: Deap of Figure 9.6 following a delete min

– **Analysis of *deap_delete_min*:**

- Time complexity of *deap_delete_min* is $O(\log n)$ as the height of the deap is $O(\log n)$.

9.3 Leftist trees

- In the preceding section we extended the definition of a priority queue by requiring that both delete max and delete min operations be permissible.
- In this section, we consider a different extension, *combine*. This requires us to **combine two priority queues into a single priority queue**.
 - One application for this is when the server for one priority queue shuts down. At this time, it is necessary to combine its priority queue with that of a functioning server.

- Simple comparisons
 - Let n be the total number of elements in the two priority queues that are to be combined.
 - If heaps are used to represent priority queues, then the combine operation takes $O(n)$ time.
 - Using a **leftist tree**, the combine operation as well as the normal priority queues operations take logarithmic time.
- In order to define a leftist tree, we need to introduce the concept of an extended binary tree.
 - An **extended binary tree** is a binary tree in which all empty binary subtrees have been replaced by a square node.
 - Figure 9.11 shows two example binary trees. Their corresponding extended binary tree are shown in Figure 9.12.
 - The square nodes in an extended binary tree are called **external nodes**. The original (circular) nodes of the binary tree are called **internal nodes**.

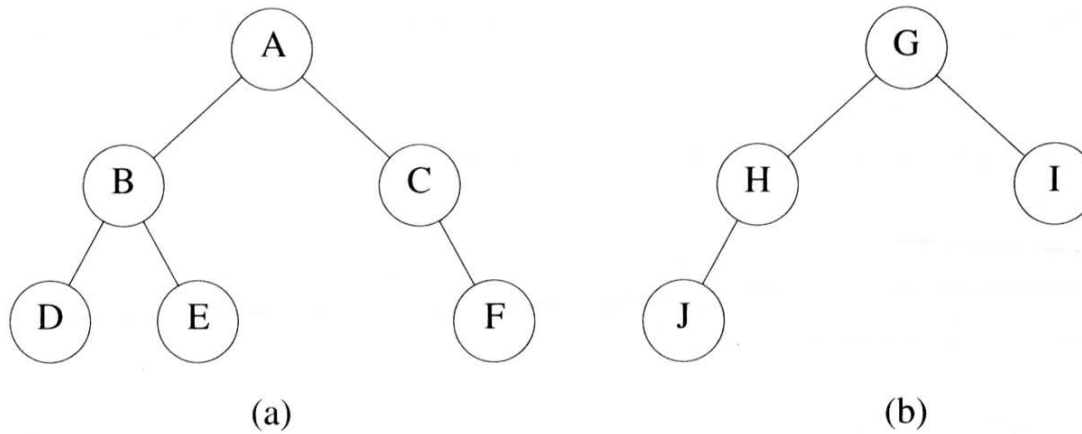


Figure 9.11: Two binary trees

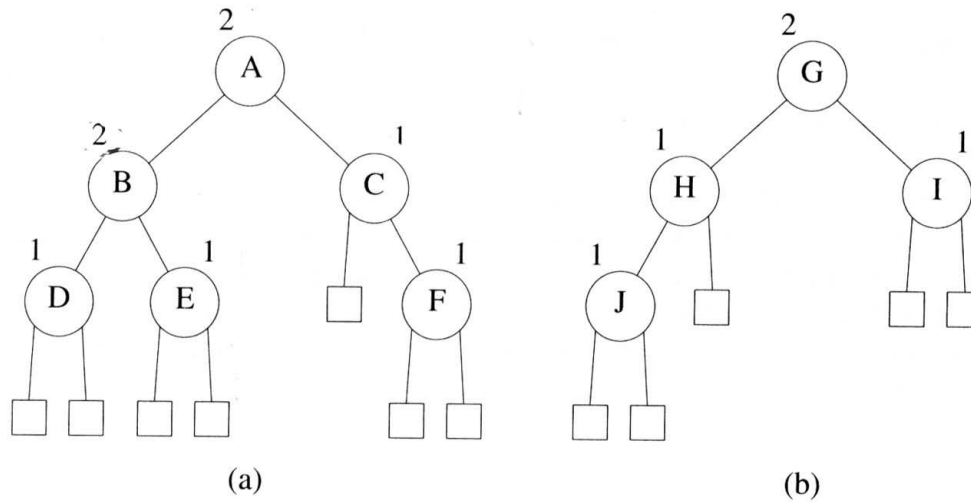


Figure 9.12: Extended binary trees corresponding to Figure 9.11

- Let x be a node in an extended binary tree. Let $left_child(x)$ and $right_child(x)$, respectively, denote the left and right children of the internal node x . Define *shortest* (x) to be the length of a shortest path from x to an external node.

$$shortest(x) = \begin{cases} 0, & \text{if } x \text{ is an external node} \\ 1 + \min\{shortest(left_child(x)), shortest(right_child(x))\}, & \text{otherwise} \end{cases}$$

– Definition:

A *leftist tree* is a binary tree such that if it is not empty, then $\text{shortest}(\text{left_child}(x)) \geq \text{shortest}(\text{right_child}(x))$ for every internal node x .

– Example:

The binary tree of Figure 9.11(a) which corresponds to the extended binary tree of Figure 9.12(a) is not a leftist tree as $\text{shortest}(\text{left_child}(C)) = 0$ while $\text{shortest}(\text{right_child}(C)) = 1$. The binary tree of Figure 9.11(b) is a leftist tree.

– Lemma 9.1:

Let x be the root of a leftist tree that has n (internal) nodes.

$$n \geq 2^{\text{shortest}(x)} - 1$$

The rightmost root to external node path is the shortest root to external node path. Its length is $\text{shortest}(x)$.

– Definition:

- A *min-leftist tree* (*max leftist tree*) is a leftist tree in which the key value in each node is no larger (smaller) than the key values in its children (if any). In other words, a min (max) leftist tree is a leftist tree that is also a min (max) tree.

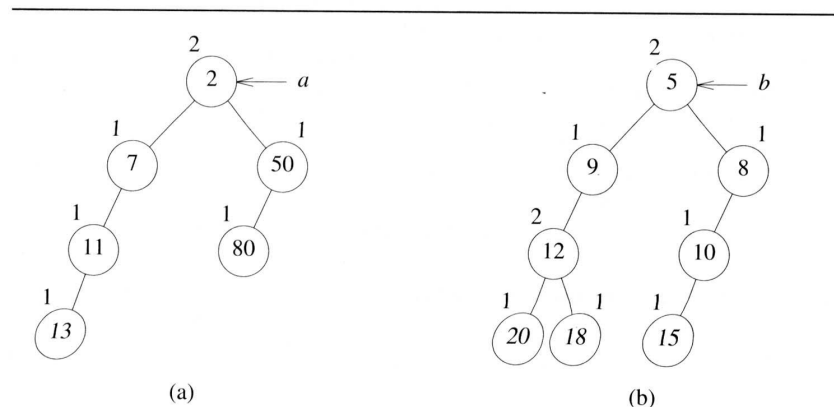
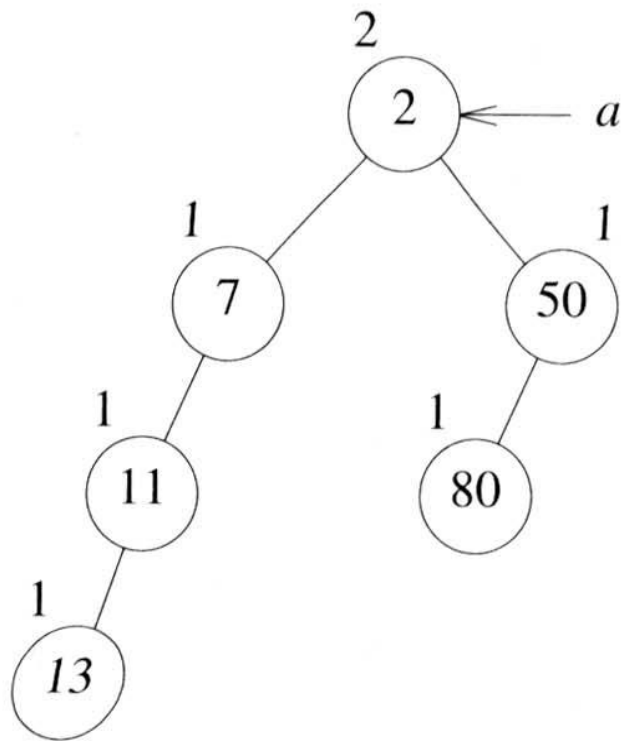


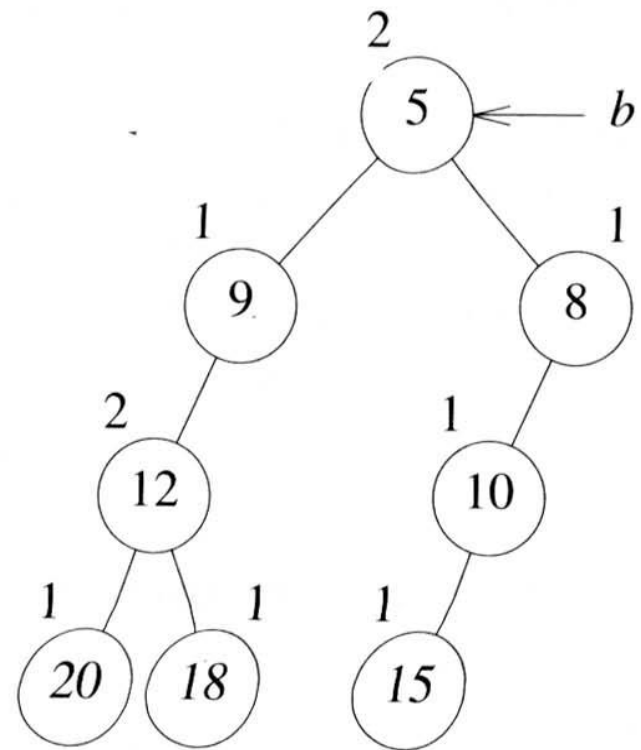
Figure 9.13: Example min leftist trees

- Both **insert** and **delete** min operations can be implemented by using the **combine** operation.
 - To insert an element, x , into a min-leftist tree, a , we first create a min-leftist tree, b , that contains the single element x . Then we combine the min-leftist trees a and b .
 - To delete the min element from a nonempty min-leftist tree, a , we combine the min-leftist trees $a->left_child$ and $a->right_child$ and delete node a .

- **Combine** the min-leftist trees a and b .
 - We obtain a new binary tree containing all elements in a and b by following the rightmost paths in a and/or b . This binary tree has the property that the key in each node is no larger than the keys in its children (if any).
 - Next, we interchange the left and right subtrees of nodes as necessary to convert this binary tree into a leftist tree.



(a)



(b)

Figure 9.13: Example min leftist trees

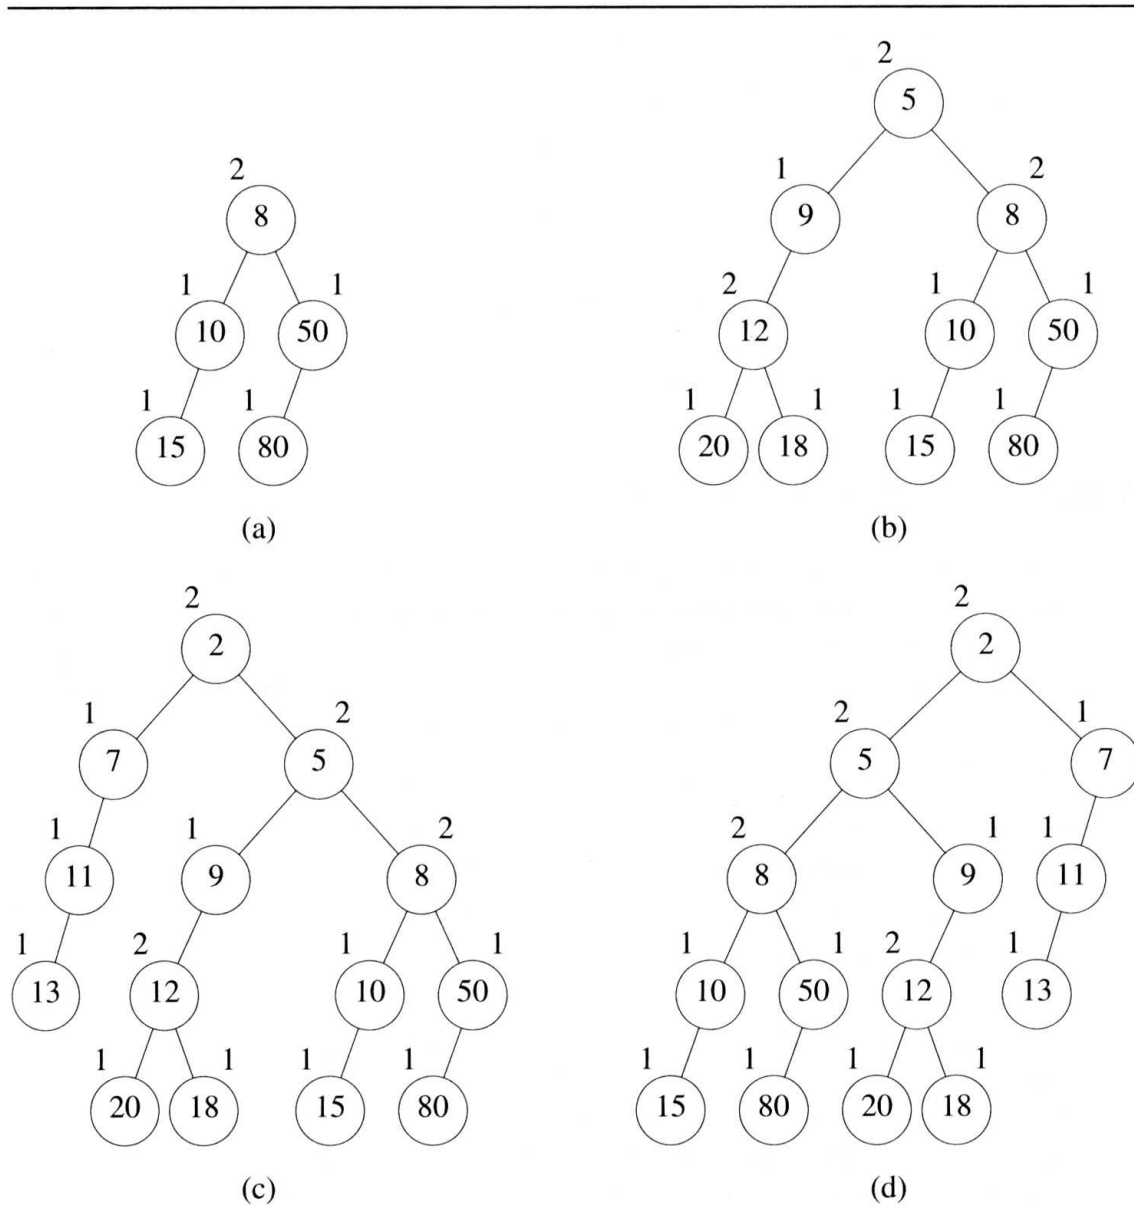


Figure 9.14: Combining the min leftist trees of Figure 9.13

– Combine方式

1. 比較兩個min-leaf tree 的root, 較小的root a 是合併後的root.
2. 合併後root a的left subtree不變.
3. Right subtree則以root a 的right subtree與另一棵tree合併而得.
4. 必要時可以對調left與right subtrees以符合leaf tree的條件.

Insert : reduced to a combine operation.

Delete : delete the root and combine the two subtrees.

9.4 Binomial heaps

- Cost amortization
 - A binomial heap is a data structure that supports the same functions as supported by leftist trees.
 - Unlike leftist trees where an individual operation can be performed in $O(\log n)$ time, if we amortize part of the cost of expensive operations over the inexpensive ones, then the amortized complexity of an individual operation is either $O(1)$ or $O(\log n)$ depending on the type of the operation.

I1	I2	D1	I3	I4	I5	I6	D2	I7
----	----	----	----	----	----	----	----	----

actual cost

1	1	8	1	1	1	1	10	1
---	---	---	---	---	---	---	----	---

Total = 25

amortized cost

2	2	6	2	2	2	2	6	1
---	---	---	---	---	---	---	---	---

Total = 25

- We can make the claim that the actual cost of any insert/delete min sequence is no more than $2*i+6*d$ where i and d are respectively, the number of insert and delete min operations in the sequence.
- Suppose that actual cost of a delete min is no more than ten, while that of an insert is one. Using actual costs, we can conclude that the sequence cost is no more than $i*10*d$.
- Combine these two bounds, we obtain $\min\{2*i+6*d, i+10*d\}$ as a bound on the sequence cost. Hence, using the notion of cost amortization it is possible to obtain tighter bounds on the complexity of a sequence of operations.

- We shall use the notion of cost amortization to show that while individual delete operations on a binomial heap may be expensive, the cost of any sequence of binomial heap operations is actually quite small.

- Definition of binomial heaps
 - As in case of heaps and leftist trees, there are two varieties of binomial heaps: min and max.
 - A *min-binomial heap* is a collection of min-trees while a *max-binomial heap* is a collection of max-trees.
 - We shall explicitly consider min-binomial heaps only. These will be referred to as *B-heaps*.

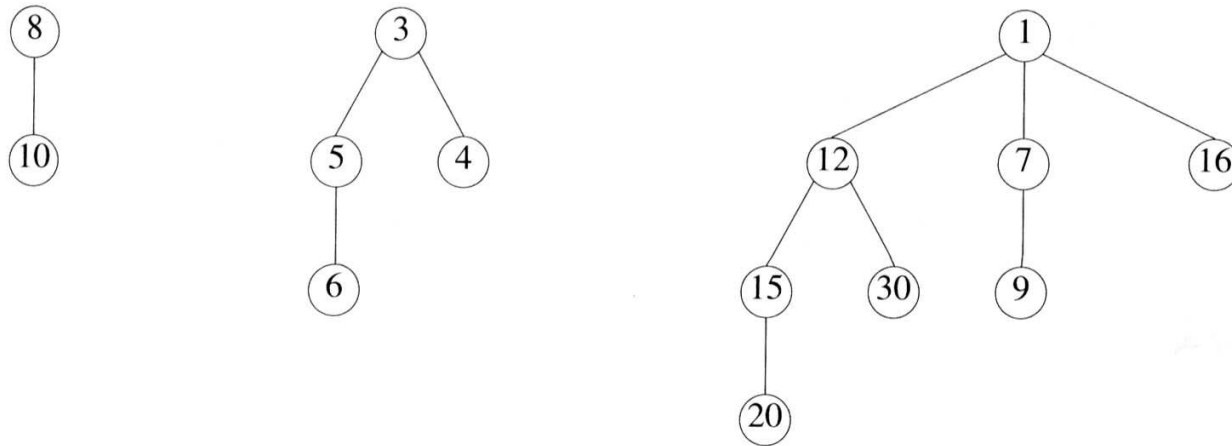


Figure 9.15: A B-heap with three min-trees

- Using B-heaps, we can perform an insert and a combine in $O(1)$ actual and amortized time and a delete min in $O(\log n)$ amortized time.
- B-heaps are represented using nodes that have the fields: *degree*, *child*, *left_link*, *right_link*, and *data*.
 - The *degree* of a node is the number of children it has.
 - The *child* field is used to point to any one of its children (if any).
 - The *left_link* and *right_link* fields are used to maintain doubly linked circular list of siblings.
 - All the children of a node form a doubly linked circular list and the node points to one of these children.
 - The roots of the min-trees that comprise a B-heap are linked to form a doubly linked circular list.
 - The B-heap is then pointed at by a single pointer to the min-tree root with smallest key.

- Figure 9.16 shows the representation for the example of Figure 9.15.

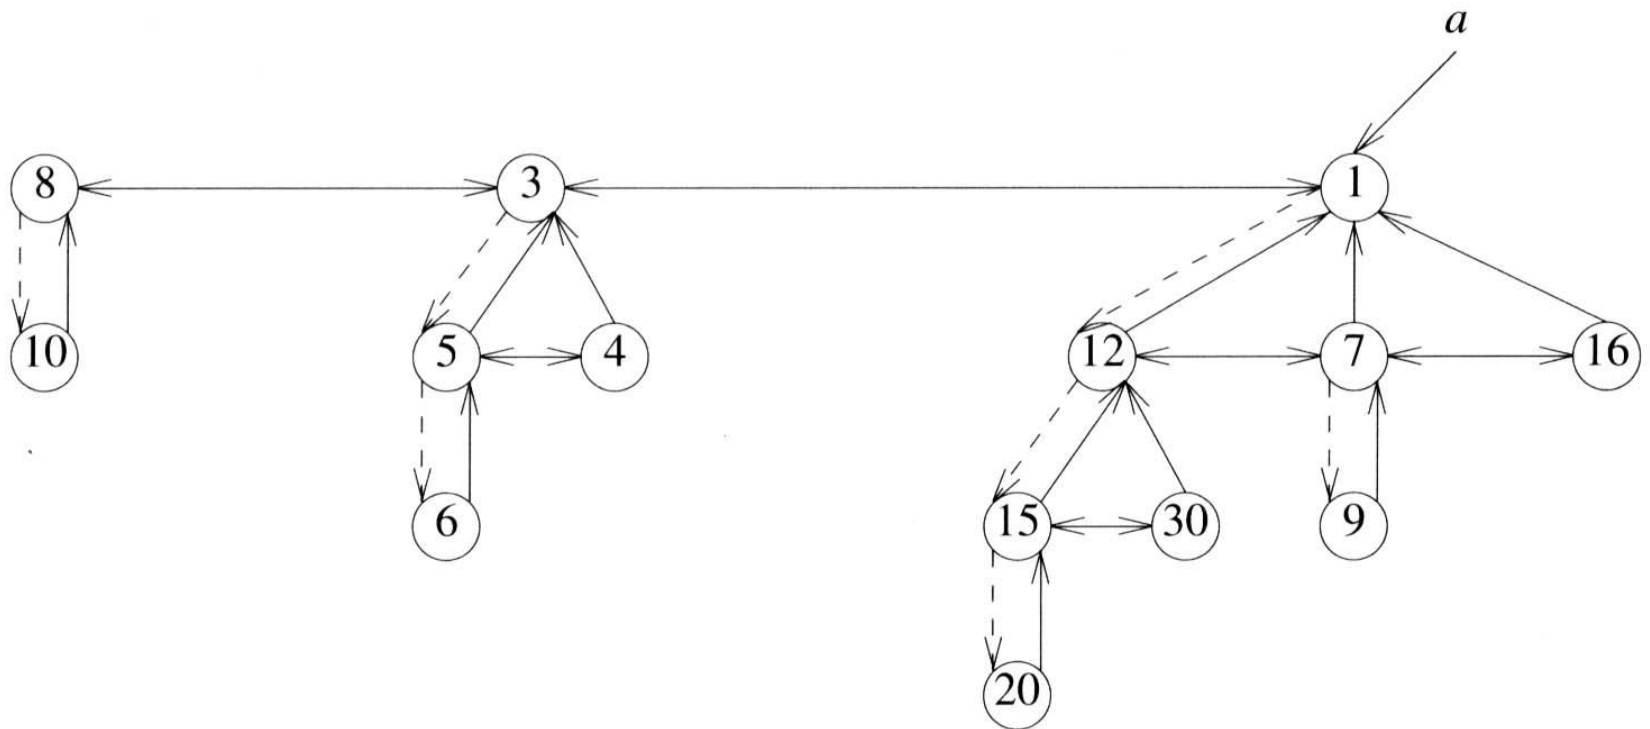


Figure 9.16: B-heap of Figure 9.15 showing parent pointers and sibling lists

- Insertion into a binomial heap
 - We insert an element, x , into an B-heap, a , by first **putting x into a new node** and placing this node into the doubly linked circular list pointed at by a .
 - We **reset a** to this new node only **if a is $NULL$ or** x 's key is smaller than the key in the node pointed at by a .
 - It is evident that we can perform these insertion steps in $O(1)$ time.

Insert 時，不作 combine 的動作！

- Combine
 - To combine two nonempty B-heaps a and b , we combine the top doubly linked circular list of a and b into single double linked circular list.
 - The new B-heap pointer is either a or b depending on which has the smaller key. This can be determined with a single comparison.
 - Since two doubly linked circular lists can be combined into a single one in $O(1)$ time, a combine takes only $O(1)$ time.

- Deletion of min element
 - Let a be the pointer of the B-heap from which the min element is to be deleted.
 - If a is *NULL*, then the B-heap is empty and a deletion cannot be performed.
 - Assume that a is not *NULL*.

a points to the node that contains the min element.

This node is deleted from its doubly linked circular list.

The new B-heap consists of the remaining min-trees and the sub min-trees of the deleted root. (Figure 9.17 shows the situation for the example of Figure 9.15)

Repeatedly join together pairs of min-trees that have the same degree. (The resulting min-tree collection is that of Figure 9.19)

Form the doubly linked circular list of min-tree roots. We also reset the B-heap pointer so as to point to the min-tree root with smallest key.

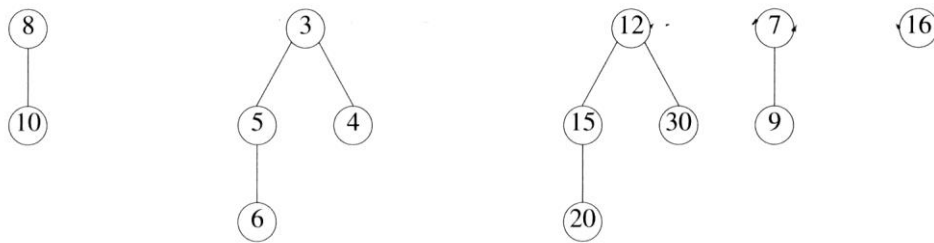


Figure 9.17: The B-heap of Figure 9.15 following the deletion of the min element

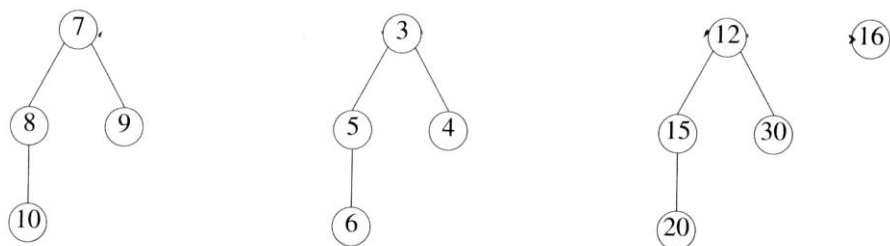


Figure 9.18: The B-heap of Figure 9.17 following the joining of the two degree one min-trees

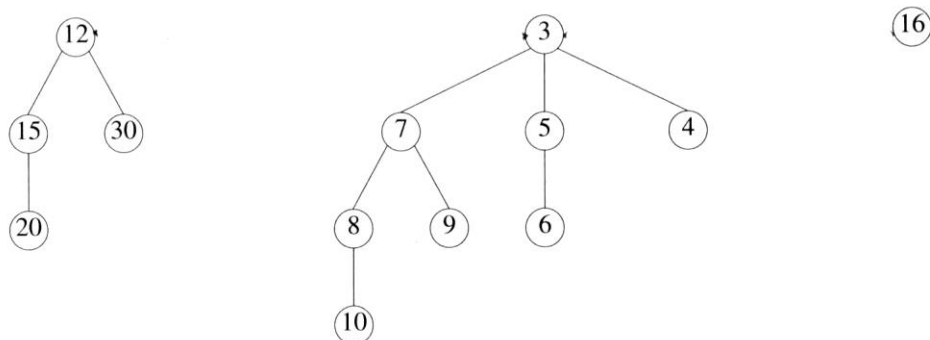


Figure 9.19: The B-heap of Figure 9.18 following the joining of two degree two min-trees

Root 的 degree:

- 是0的，有 2^0 個nodes (B_0)
- 是1的，有 2^1 個nodes (B_1)
- 是2的，有 2^2 個nodes (B_2)
- 是3的，有 2^3 個nodes (B_3)

整個 B-heap 就成為

$$a_0B_0 + a_1B_1 + a_2B_2 + a_3B_3 \dots$$

- How to repeatedly join together pairs of min-trees that have the same degree?
 - The degree of a nonempty min-tree is the degree of its root.
 - This min-tree joining is done by making the min-tree whose root has larger key a subtree of the other.
 - When two min-trees are joined, the degree of the resulting min-tree is one larger than the original degree of each min-tree and the number of min-trees decreases by one.
 - Example: We may first join the min-trees with roots 8 and 7. The min-tree with root 8 is made a subtree of the min-tree with root 7. We now have the min-tree collection of Figure 9.18.

{Delete the min element from a B-heap a , this element is returned in x }

Step 1: [Handle empty B-heap] if ($a = \text{NULL}$) *deletion – error* else perform steps 2 - 4;

Step 2: [Deletion from nonempty B-heap] $x = a \rightarrow \text{data}$; $y = a \rightarrow \text{child}$; Delete a from its doubly linked circular list; Now, a points to any remaining node in the resulting list; If there is no such node, then $a = \text{NULL}$;

Step 3: [Min-tree joining] Consider the min-trees in the lists a and y ; Join together pairs of min-trees of the same degree until all remaining min-trees have different degree;

Step 4: [Form min-tree root list] Link the roots of the remaining min-trees (if any) together to form a doubly linked circular list; Set a to point to the root (if any) with minimum key;

- Analysis

- **Definition**

The *binomial tree*, B_k , of *degree* k is a tree such that if $k = 0$, then the tree has exactly one node and if $k > 0$, then it consists of a root whose **degree is k** and whose **subtree are B_0, B_1, \dots, B_{k-1}** .

- Example: The min-trees of Figure 9.15 are B_1, B_2 , and B_3 , respectively.

- One may verify that **B_k has exactly 2^k nodes**.
 - Further, if we start with a collection of empty B-heaps and perform only the operations insert, combine, and delete min, then the min-trees in each B-heap are binomial trees.

– **Lemma 9.2:**

Let a be a B-heap with n elements that results from a sequence of insert, combine, and delete min operations performed on initially empty B-heaps. **Each min-tree in a has degree $\leq \log_2 n$.** Consequently, $MAX_DEGREE \leq \lfloor \log_2 n \rfloor$ and the actual **cost of a delete min** is

$O(\log_2 n + s)$,

where **s** be the number of min-trees in a and y ($y = a \rightarrow child$).

– **Theorem 9.1:**

If a sequence of n insert, combine, and delete min operations is performed on initial empty B-heaps, then we can amortize costs such that the amortized time complexity of each insert and **combine** is $O(1)$ and that of each **delete min** is $O(\log n)$.

9.5 Fibonacci heaps

- Definition (e.g., 1, 1, 2, 3, 5, 8, 13, 21.....)
 - A **Fibonacci heap** is a data structure that supports the three binomial heap operations: **insert**, **delete min**, and **combine** as well as the operations:
 1. *delete*, delete the element in a specified node (done in $O(1)$ amortized time)
 2. *decrease key*, decrease the key of a specified node by a given positive amount (in $O(\log n)$ amortized time).
 - The binomial heap operations can be performed in the same asymptotic times using a Fibonacci heap as using a binomial heap. (作動作的時間 **F-heap = B-heap**)

- There are two varieties of Fibonacci heaps: min and max. A *min-Fibonacci heap* is a collection of min-trees while a *max-Fibonacci heap* is a collection of max-trees.
 - We shall explicitly **consider min-Fibonacci heaps only**.
 - These will be referred to as *F-heaps*.
- B-heaps are a special case of F-heaps.
- To represent an F-heap, the B-heap representation is augmented by adding two fields: *parents* and *child_cut* to each node.
 - The *parents* field is used to point to the node's parent.
 - The significance of the *child_cut* field will be described later.
- The basic operations: insert, delete min, and combine are performed exactly as for the case of B-heaps.

- Deletion from an F-heap
 - To delete an arbitrary node b from the F-heap a , we do the following:
 - If $a = b$, then do a delete min; otherwise do steps 2, 3, and 4 below.
 - Delete b from the doubly linked list it is in.
 - Combine the doubly linked list of b 's children with the doubly linked list of a 's min-tree roots to get a single doubly linked list. Trees of equal degree are not joined together as in a delete min.
 - Dispose of node b .

- Example: if we delete the node containing 12 from the F-heap of Figure 9.15, we get the F-heap of Figure 9.20. The actual cost of an arbitrary delete is $O(1)$ unless the min element is being deleted. In this case the deletion time is the time for a delete min operation.

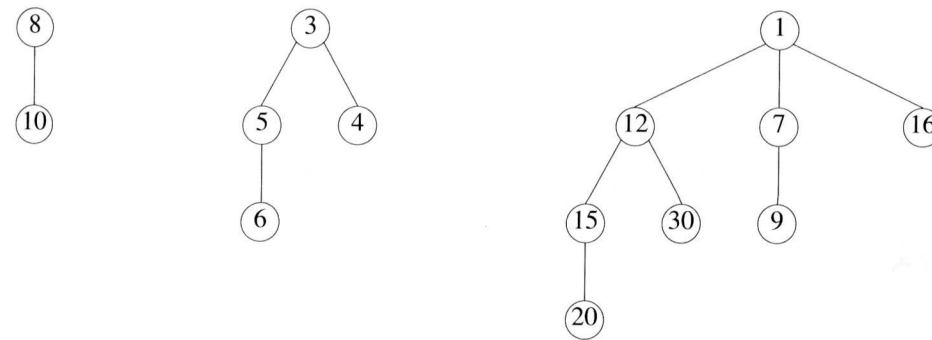


Figure 9.15: A B-heap with three min-trees

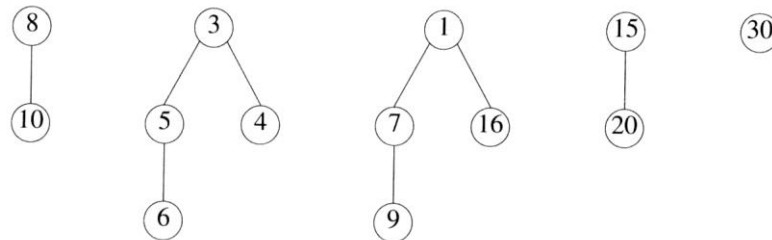


Figure 9.20: F-heap of Figure 9.15 following the deletion of 12

- Decrease key
 - To decrease the key in node b we do the following:
 - Reduce the key in b .
 - If b is not a min-tree root and its key is smaller than that in its parent, then delete b from its doubly linked list and insert it into the doubly linked list of min-tree roots.
 - Change a to point to b in case the key in b is smaller than that in a .

- Example: suppose we decrease the key 15 in the F-heap of Figure 9.15 by 4. The resulting F-heap is shown in Figure 9.21. The cost of performing a decrease key is $O(1)$.

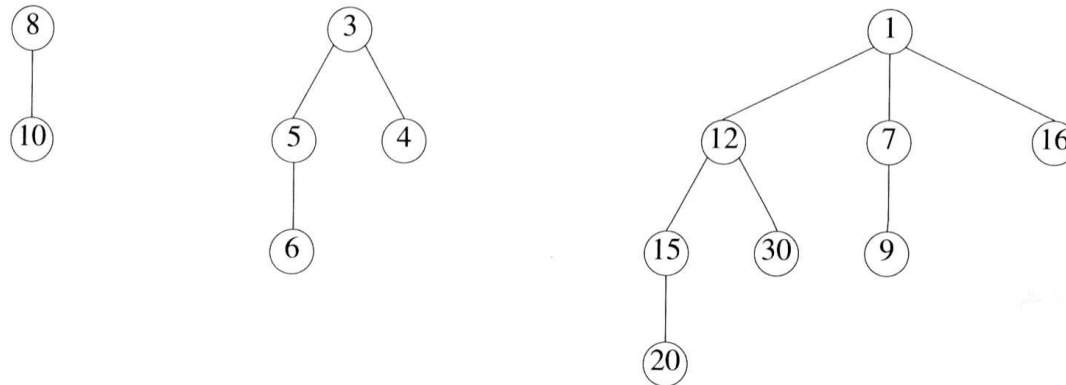


Figure 9.15: A B-heap with three min-trees

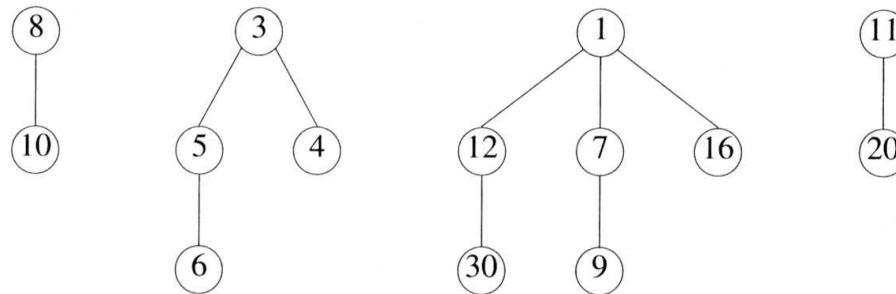


Figure 9.21: F-heap of Figure 9.15 following the reduction of 15 by 4

- If an F-heap is used to represent S' , the complexity of the shortest path algorithm becomes $O(n \log n + e)$. (e : the number of edges)
 - This is an asymptotic improvement over the implementation discussed in Chapter 6 if the graph does not have $\Omega(n^2)$ edges.
- If this single source algorithm is used n times, once with each of the n vertices in the graph as the source, then we can find a shortest path between every pair of vertices in $O(n^2 \log n + ne)$.
 - Once again, this represents an asymptotic improvement over the $O(n^3)$ dynamic programming algorithm of Chapter 6 for graphs that do not have $\Omega(n^2)$ edges.
- It is interesting to note that $O(n \log n + e)$ is the best possible implementation of the single source algorithm of Chapter 6 as the algorithm must examine each edge and may be used to sort n numbers (which requires $O(n \log n)$ time).