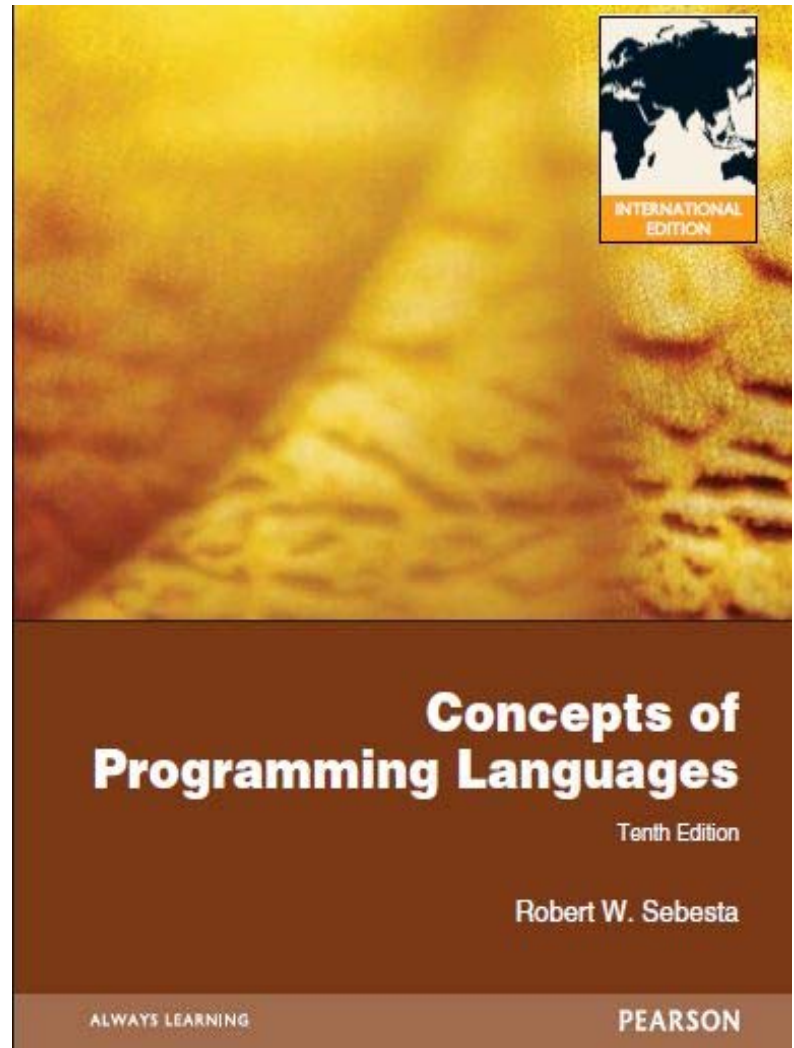


Programming Language

Instructor:

Min-Chun Hu

anita_hu@mail.ncku.edu.tw



Lecture 8 Concurrency

- ❑ Introduction
- ❑ Introduction to Subprogram-Level Concurrency
- ❑ Semaphores
- ❑ Monitors
- ❑ Message Passing
- ❑ Ada Support for Concurrency
- ❑ Java Threads
- ❑ C# Threads
- ❑ Concurrency in Functional Languages
- ❑ Statement-Level Concurrency

Introduction

- Concurrency can occur at four levels:
 - ▣ Machine instruction level
 - ▣ High-level language statement level
 - ▣ Unit level
 - ▣ Program level Concurrent control methods increase programming flexibility
- A **concurrent** algorithm is **scalable** if the speed of its execution increase when more processors are available. ex: sort, search
 - Instruction level is the execution of two or more machine instructions simultaneously
 - Statement level is the execution of two or more statements simultaneously
 - Unit level is the execution of two or more subprogram units simultaneously
 - Program level is the execution of two or more programs simultaneously

Multiprocessor Architectures

- Late 1950s – one general-purpose processor and one or more special-purpose processors for input and output operations
- Early 1960s – multiple complete processors, used for program-level concurrency
- Mid-1960s – multiple partial processors, used for instruction-level concurrency

並行性（concurrency）是指在一個系統中，擁有多個計算，這些計算有同時執行的特性，而且他們之間有著潛在的互動。因此系統可進行的執行路徑會有相當多個，而且結果可能具有不確定性。並行計算可能會在具備多核心的同一個晶片中複合執行，以優先分時執行緒在同一個處理器中執行，或在不同的處理器執行。

Multiprocessor Architectures (Cont.)

- Single-Instruction Multiple-Data (SIMD) machines
 - ▣ Multiple processors that execute the same instruction simultaneously on different data
 - ▣ Applications: graphics/video processing
- Multiple-Instruction Multiple-Data (MIMD) machines
 - ▣ Multiple processors that operate independently but whose operations can be synchronized.
- A primary focus of this chapter is shared memory MIMD machines (multiprocessors)

Categories of Concurrency

- Categories of Concurrency:
 - ▣ *Physical concurrency* – Multiple independent processors (multiple threads of control)
 - ▣ *Logical concurrency* – The appearance of physical concurrency is presented by time-sharing one processor (interleaving)
- A *thread of control* in a program is the sequence of program points reached as control flows through the program
- *Coroutines (quasi-concurrency) have a single thread of control*

執行緒（Thread）是作業系統能夠進行運算排程的最小單位。它被包含在行程之中，是行程中的實際運作單位。一條執行緒指的是行程中一個單一順序的控制流，一個行程中可以並行多個執行緒，每條執行緒並列執行不同的任務。

Motivations for the Use of Concurrency

- Multiprocessor computers capable of physical concurrency are now widely used
- Even if a machine has just one processor, a program written to use concurrent execution can be faster than the same program written for nonconcurrent execution
- Many real-world situations involve concurrency and many program applications are now spread over multiple machines, either locally or over a network

Introduction to Subprogram-Level Concurrency

- A *task* or *process* or *thread* is a program unit that can be in concurrent execution with other program units
- Tasks differ from ordinary subprograms in that:
 - ▣ A task may be implicitly started
 - ▣ When a program unit starts the execution of a task, it is not necessarily suspended
 - ▣ When a task's execution is completed, control may not return to the caller
- Tasks usually work together

Two General Categories of Tasks

- *Heavyweight tasks* execute in their own address space 程序 (process)
- *Lightweight tasks* all run in the same address space – more efficient 執行緒 (thread)
- A task is *disjoint* if it does not communicate with or affect the execution of any other task in the program in any way

Task Synchronization

- A mechanism that controls the order in which tasks execute
- Two kinds of synchronization
 - ▣ Cooperation synchronization
 - ▣ Competition synchronization
- Task communication is necessary for synchronization, provided by:
 - Shared nonlocal variables
 - Parameters
 - Message passing
- Cooperation synchronization is required between two tasks that when the second task must wait for the first task to finish executing before it may proceed.
- Competition synchronization is required between two tasks when both require the use of the same resource that cannot be simultaneously used.

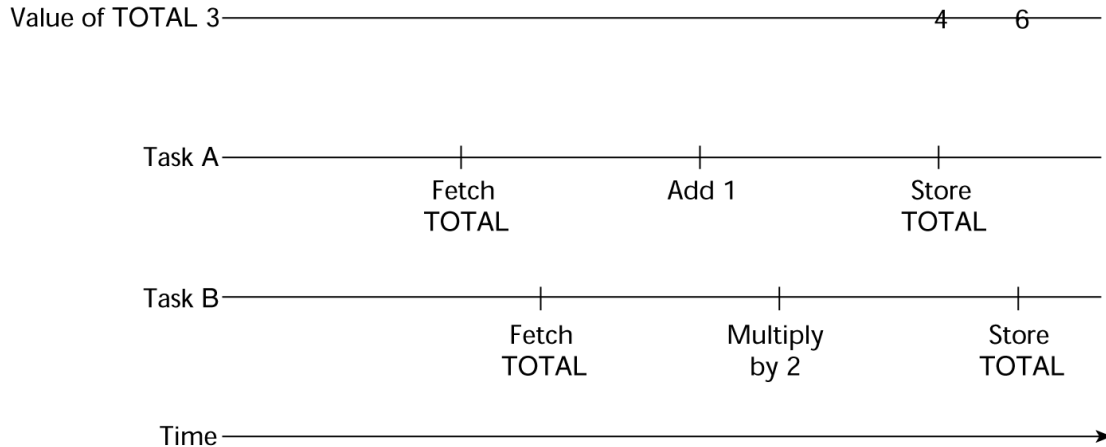
Kinds of synchronization

- *Cooperation:*
 - ▣ Task A must wait for task B to complete some specific activity before task A can continue its execution,
 - ▣ e.g., the producer–consumer problem
- *Competition:*
 - ▣ Two or more tasks must use some resource that cannot be simultaneously used, e.g., a shared counter
 - ▣ Competition is usually provided by mutually exclusive access (approaches are discussed later)

Need for Competition Synchronization

Task A: $TOTAL = TOTAL + 1$

Task B: $TOTAL = 2 * TOTAL$



- Depending on order, there could be four different results

Race condition

Scheduler

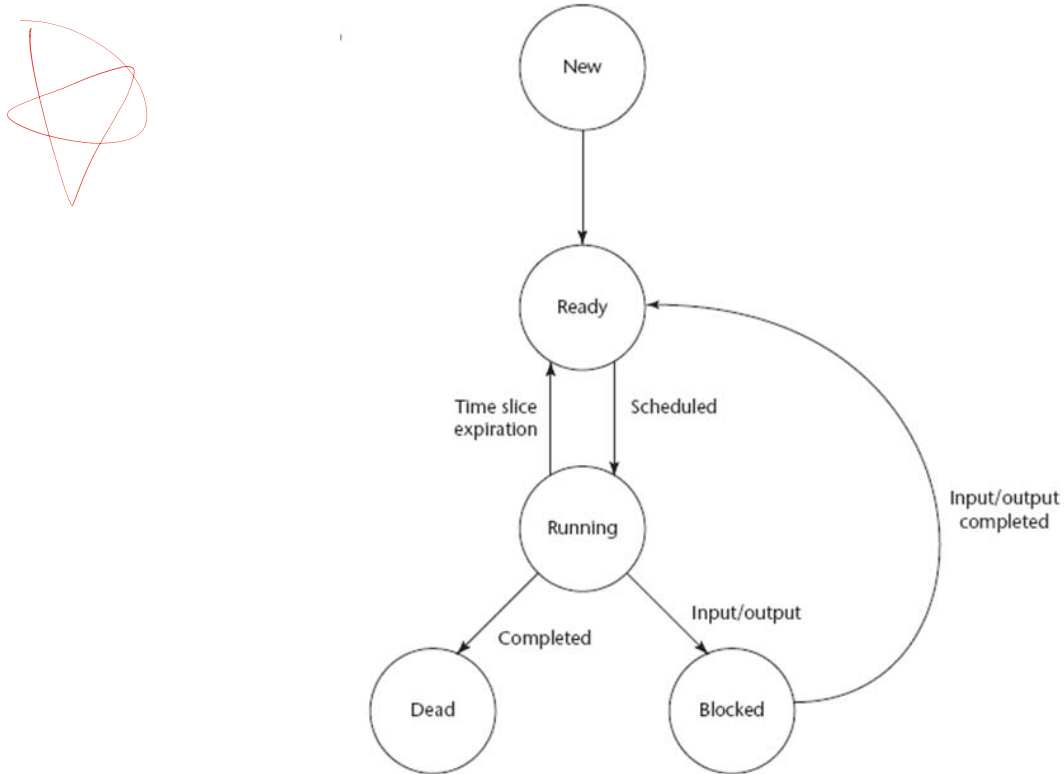
- Providing synchronization requires a mechanism for **delaying task execution**
- Task execution control is maintained by a program called the **scheduler**, which maps task execution onto available processors

排程（scheduler），是將任務分配至資源的過程，排班首要面對的就是效率問題。

Task Execution States

- *New* – created but not yet started
- *Ready* – ready to run but not currently running (no available processor)
 - ▣ Tasks ready to run are stored in the *task ready queue*
- *Running* – currently executing
- *Blocked* – has been running, but cannot now continue (usually waiting for some event to occur)
- *Dead* – no longer active in any sense (execution completed or explicitly killed by the program)

Flow Diagram of Task States



Liveness and Deadlock

- *Liveness* is a characteristic that a program unit may or may not have
 - ▣ In sequential code, it means the unit will eventually complete its execution
 - ▣ In a concurrent environment, a task can easily lose its liveness
 - ▣ If all tasks in a concurrent environment lose their liveness, it is called *deadlock*

If two tasks have a resource that the other needs and neither releases the resource it possesses this guarantees that neither program will ever complete normally.

Design Issues for Concurrency

- Competition and cooperation synchronization
 - ▣ The most important issue
- Controlling task scheduling
- How can an application influence task scheduling
- How and when tasks start and end execution
- How and when are tasks created

Methods of Providing Synchronization

- Three methods proving mutually exclusive access:
 - ▣ Semaphores
 - ▣ Monitors
 - ▣ Message Passing

Semaphores

- Dijkstra – 1965
- A *semaphore* is a data structure consisting of a counter and a queue for storing task descriptors
 - ▣ A task descriptor is a data structure that stores all of the relevant information about the execution state of the task
- Semaphores can be used to implement guards on the code that accesses shared data structures
- Semaphores have only two operations, *wait* and *release* (originally called *P* and *V* by Dijkstra)
- Semaphores can be used to provide both competition and cooperation synchronization

Counting Semaphore

Binary Semaphore

- A very simple mechanism used to provide synchronization of tasks.
- A semaphore provides limited access to a data structure by placing guards around the code that accesses the structure.

Cooperation Synchronization with Semaphores

- Example: A shared buffer
- The buffer is implemented as an ADT with the operations `DEPOSIT` and `FETCH` as the only ways to access the buffer
存 取
- Use two semaphores for cooperation: `emptyspots` and `fullspots`
- The semaphore counters are used to store the numbers of empty spots and full spots in the buffer

Cooperation Synchronization with Semaphores (Cont.)

- **DEPOSIT must first check `emptyspots` to see if there is room in the buffer**
 - ▣ If there is room, the counter of `emptyspots` is decremented and the value is inserted
 - ▣ If there is no room, the caller is stored in the queue of `emptyspots`
 - ▣ When **DEPOSIT** is finished, it must increment the counter of `fullspots`

Cooperation Synchronization with Semaphores (Cont.)

- `FETCH` must first check `fullspots` to see if there is a value
 - ▣ If there is a full spot, the counter of `fullspots` is decremented and the value is removed
 - ▣ If there are no values in the buffer, the caller must be placed in the queue of `fullspots`
 - ▣ When `FETCH` is finished, it increments the counter of `emptyspots`
- The operations of `FETCH` and `DEPOSIT` on the semaphores are accomplished through two semaphore operations named *wait* and *release*

Semaphores: Wait and Release Operations

```
wait(aSemaphore)
if aSemaphore's counter > 0 then
    decrement aSemaphore's counter
else
    put the caller in aSemaphore's queue
    attempt to transfer control to a ready task
    -- if the task ready queue is empty,
    -- deadlock occurs
end

release(aSemaphore)
if aSemaphore's queue is empty then
    increment aSemaphore's counter
else
    put the calling task in the task ready queue
    transfer control to a task from aSemaphore's queue
end
```

Producer and Consumer Tasks

```
semaphore fullspots, emptyspots;
fullspots .count = 0;
emptyspots.count = BUFLen;
task producer;
    loop
        -- produce VALUE --
        wait (emptyspots); {wait for space}
        DEPOSIT(VALUE);
        release(fullspots); {increase filled}
    end loop;
end producer;
task consumer;
    loop
        wait (fullspots); {wait till not empty}
        FETCH(VALUE);
        release(emptyspots); {increase empty}
        -- consume VALUE --
    end loop;
end consumer;
```


Competition Synchronization with Semaphores

- A third semaphore, named `access`, is used to control access (competition synchronization)
 - ▣ The counter of `access` will only have the values 0 and 1
 - ▣ Such a semaphore is called a *binary semaphore*

Producer Code for Semaphores

```
semaphore access, fullspots, emptyspots;
access.count = 0;
fullspots    .count = 0;
emptyspots.count = BUFLen;
task producer;
    loop
        -- produce VALUE --
        wait(emptyspots); {wait for space}
        wait(access);     {wait for access}
        DEPOSIT(VALUE);
        release(access);  {relinquish access}
        release(fullspots); {increase filled}
    end loop;
end producer;
```

Consumer Code for Semaphores

```
task consumer;  
  loop  
    wait(fullspots); {wait till not empty}  
    wait(access);    {wait for access}  
    FETCH(VALUE);  
    release(access); {relinquish access}  
    release(emptyspots); {increase empty}  
    -- consume VALUE --  
  end loop;  
end consumer;
```

Evaluation of Semaphores

- Misuse of semaphores can cause failures in cooperation synchronization, e.g., the buffer will **overflow** if the wait of **fullspots** is left out
underflow
emptyspot
- Misuse of semaphores can cause failures in competition synchronization, e.g., the program will **deadlock** if the release of access is left out

Monitors

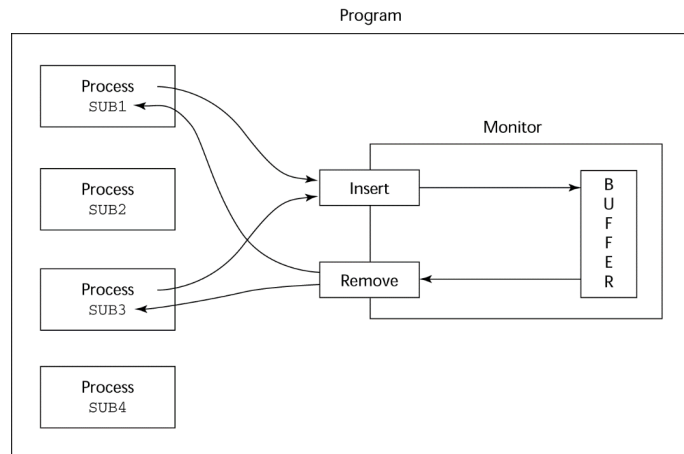
- Ada, Java, C#
- The idea: encapsulate the shared data and its operations to restrict access
- A monitor is an abstract data type for shared data

Competition Synchronization

- Shared data is resident in the monitor (rather than in the client units)
- All access resident in the monitor
 - ▣ Monitor implementation guarantee synchronized access by allowing only one access at a time 具Mutually exclusive 的機制
 - ▣ Calls to monitor procedures are implicitly queued if the monitor is busy at the time of the call

Cooperation Synchronization

- Cooperation between processes is still a programming task
 - ▣ Programmer must guarantee that a shared buffer does not experience underflow or overflow



Evaluation of Monitors

- A better way to provide competition synchronization than are semaphores
- Semaphores can be used to implement monitors
- Monitors can be used to implement semaphores
- Support for cooperation synchronization is very similar as with semaphores, so it has the same problems

Message Passing

- Message passing is a general model for concurrency
 - ▣ It can model both semaphores and monitors
 - ▣ It is not just for competition synchronization
- Central idea: **task communication is like seeing a doctor**--most of the time she waits for you or you wait for her, but when you are both ready, you get together, or *rendezvous*
- Message passing can be synchronous or asynchronous.
- Cooperation and competition synchronization can be conveniently handled with message passing.

Message Passing Rendezvous

- To support concurrent tasks with message passing, a language needs:
 - A mechanism to allow a task to indicate when it is willing to accept messages
 - A way to remember who is waiting to have its message accepted and some “fair” way of choosing the next message
- When a sender task’s message is accepted by a receiver task, the actual message transmission is called a *rendezvous*
匯合

Ada Support for Concurrency

- The Ada 83 Message-Passing Model
 - ▣ Ada tasks have specification and body parts, like packages; the spec has the interface, which is the collection of entry points:

```
task Task_Example is  
    entry ENTRY_1 (Item : in Integer);  
end Task_Example;
```

Task Body

- The `body` task describes the action that takes place when a rendezvous occurs
- A task that sends a message is suspended while waiting for the message to be accepted and during the rendezvous
- Entry points in the spec are described with `accept` clauses in the body

```
accept entry_name (formal parameters) do
```

```
    ...
```

藉由send message和receive message 機制保護data

```
end entry_name;
```

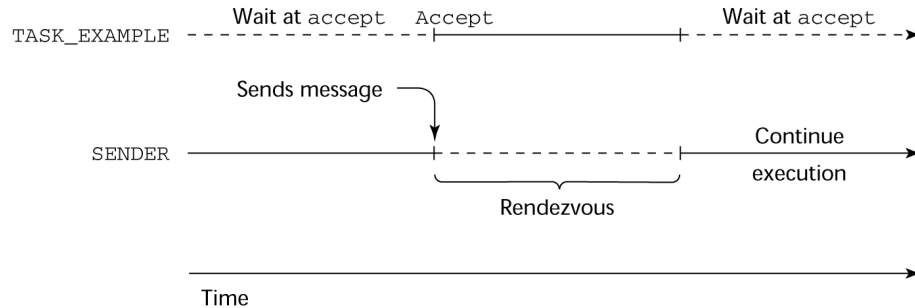
Example of a Task Body

```
task body Task_Example is
  begin
    loop
      accept Entry_1 (Item: in Float) do
        ...
      end Entry_1;
    end loop;
  end Task_Example;
```


Ada Message Passing Semantics

- The task executes to the top of the `accept` clause and waits for a message
- During execution of the `accept` clause, the sender is suspended
- `accept` parameters can transmit information in either or both directions
- Every `accept` clause has an associated queue to store waiting messages

Rendezvous Time Lines



(a) TASK_EXAMPLE waits for SENDER

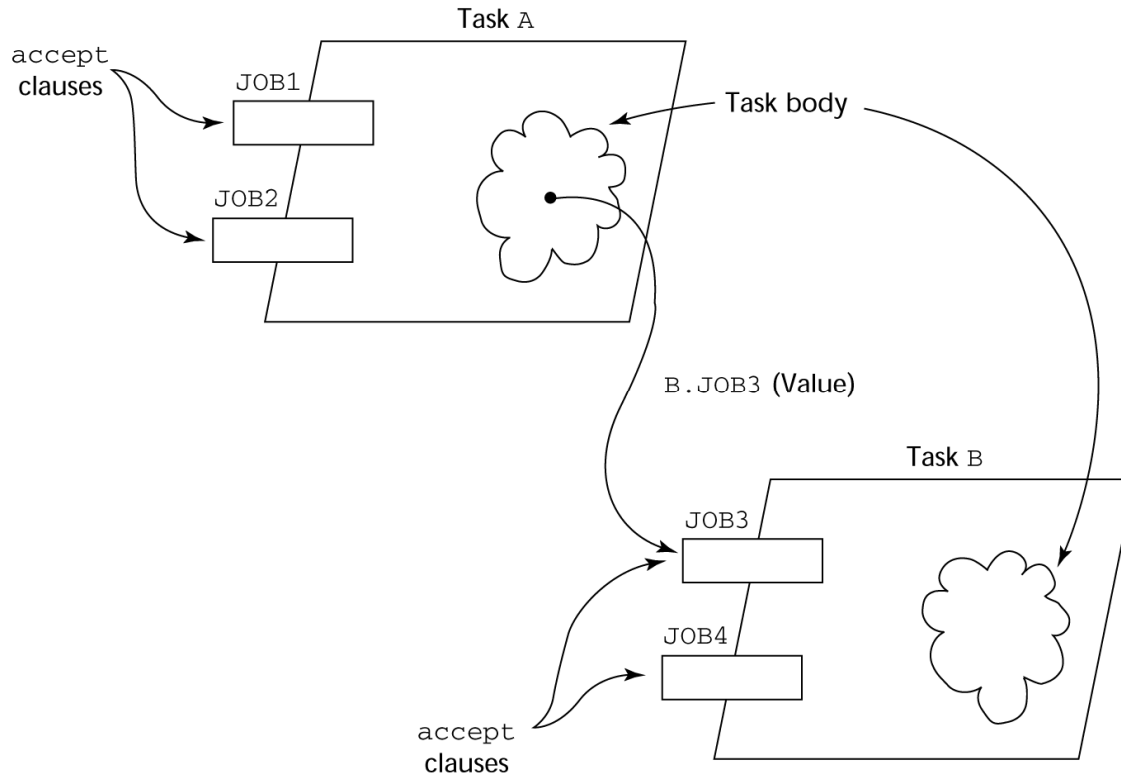


(b) SENDER waits for TASK_EXAMPLE

Message Passing: Server/Actor Tasks

- A task that has **accept** clauses, but no other code is called a *server task* (the example above is a server task)
- A task without **accept** clauses is called an *actor task*
 - ▣ An actor task can send messages to other tasks
 - ▣ Note: A sender must know the **entry** name of the receiver, but not vice versa (asymmetric)

Graphical Representation of a Rendezvous



Multiple Entry Points

- Tasks can have more than one `entry` point
 - ▣ The specification task has an `entry` clause for each
 - ▣ The task body has an `accept` clause for each `entry` clause, placed in a `select` clause, which is in a loop

A Task with Multiple Entries

```
task body Teller is
  loop
    select
      accept Drive_Up(formal params) do
        ...
      end Drive_Up;
      ...
    or
      accept Walk_Up(formal params) do
        ...
      end Walk_Up;
      ...
    end select;
  end loop;
end Teller;
```

Semantics of Tasks with Multiple `accept` Clauses

- If exactly one `entry` queue is nonempty, choose a message from it
- If more than one `entry` queue is nonempty, choose one, nondeterministically, from which to accept a message
- If all are empty, wait
- The construct is often called a `selective wait`
- `Extended accept clause` – code following the clause, but before the next clause
 - Executed concurrently with the caller

Cooperation Synchronization with Message Passing

- Provided by **Guarded accept** clauses

```
when not Full(Buffer) =>  
  accept Deposit (New_Value) do  
    ...  
  end
```

- An **accept** clause with a **when** clause is either *open* or *closed*
 - A clause whose guard is true is called *open*
 - A clause whose guard is false is called *closed*
 - A clause without a guard is always open

Semantics of `select` with Guarded `accept` Clauses:

- `select` first checks the guards on all clauses
- If exactly one is open, its queue is checked for messages
- If more than one are open, non-deterministically choose a queue among them to check for messages
- If all are closed, it is a runtime error
- A `select` clause can include an `else` clause to avoid the error
 - ▣ When the `else` clause completes, the loop repeats

Competition Synchronization with Message Passing

- Modeling mutually exclusive access to shared data
- Example--a shared buffer
- Encapsulate the buffer and its operations in a task
- Competition synchronization is implicit in the semantics of `accept` clauses
 - ▣ Only one `accept` clause in a task can be active at any given time

Partial Shared Buffer Code

```
task body Buf_Task is
  Bufsize : constant Integer := 100;
  Buf : array (1..Bufsize) of Integer;
  Filled : Integer range 0..Bufsize := 0;
  Next_In, Next_Out : Integer range 1..Bufsize := 1;
begin
  loop
    select
      when Filled < Bufsize =>
        accept Deposit(Item : in Integer) do
          Buf(Next_In) := Item;
        end Deposit;
        Next_In := (Next_In mod Bufsize) + 1;
        Filled := Filled + 1;
      or
      ...
    end loop;
end Buf_Task;
```


A Consumer Task

```
task Consumer;  
task body Consumer is  
    Stored_Value : Integer;  
    begin  
        loop  
            Buf_Task.Fetch(Stored_Value);  
            -- consume Stored_Value -  
        end loop;  
    end Consumer;
```

Task Termination

- The execution of a task is *completed* if control has reached the end of its code body
- If a task has created no *dependent tasks* and is completed, it is *terminated*
- If a task has created dependent tasks and is completed, it is not terminated until all its dependent tasks are terminated

The `terminate` Clause

- A `terminate` clause in a `select` is just a `terminate` statement
- A `terminate` clause is selected when no `accept` clause is open
- When a `terminate` is selected in a task, the task is terminated only when its master and all of the dependents of its master are either completed or are waiting at a `terminate`
- A block or subprogram is not left until all of its dependent tasks are terminated

Message Passing Priorities

- The priority of any task can be set with the `pragma Priority`
`pragma Priority (static expression);`
- The priority of a task applies to it only when it is in the task ready queue

Concurrency in Ada 95

- Ada 95 includes Ada 83 features for concurrency, plus two new features
 - ▣ Protected objects: A more efficient way of implementing shared data to allow access to a shared data structure to be done without rendezvous
 - ▣ Asynchronous communication

Ada 95: Protected Objects

- A *protected object* is similar to an abstract data type
- Access to a protected object is either through messages passed to entries, as with a task, or through protected subprograms
- A *protected procedure* provides mutually exclusive read–write access to protected objects
- A *protected function* provides concurrent read–only access to protected objects

Evaluation of the Ada

- Message passing model of concurrency is powerful and general
- Protected objects are a better way to provide synchronized shared data
- In the absence of distributed processors, the choice between monitors and tasks with message passing is somewhat a matter of taste
- For distributed systems, message passing is a better model for concurrency

Java Threads

- The concurrent units in Java are methods named `run`
 - ▣ A `run` method code can be in concurrent execution with other such methods
 - ▣ The process in which the `run` methods execute is called a *thread*

```
class myThread extends Thread
    public void run () {...}
}
```

...

```
Thread myTh = new MyThread ();
myTh.start();
```


Controlling Thread Execution

- The `Thread` class has several methods to control the execution of threads
 - ▣ The `yield` is a request from the running thread to voluntarily surrender the processor
 - ▣ The `sleep` method can be used by the caller of the method to block the thread
 - ▣ The `join` method is used to force a method to delay its execution until the run method of another thread has completed its execution

Thread Priorities

- A thread's default priority is the same as the thread that create it
 - ▣ If `main` creates a thread, its default priority is `NORM_PRIORITY`
- Threads defined two other priority constants, `MAX_PRIORITY` and `MIN_PRIORITY`
- The priority of a thread can be changed with the methods `setPriority`

Competition Synchronization with Java Threads

- A method that includes the `synchronized` modifier disallows any other method from running on the object while it is in execution

```
...  
public synchronized void deposit( int i) {...}  
public synchronized int fetch() {...}
```

- The above two methods are synchronized which prevents them from interfering with each other
- If only a part of a method must be run without interference, it can be synchronized thru `synchronized statement`
`synchronized (expression)`
statement

Cooperation Synchronization with Java Threads

- Cooperation synchronization in Java is achieved via `wait`, `notify`, and `notifyAll` methods
 - ▣ All methods are defined in `Object`, which is the root class in Java, so all objects inherit them
- The `wait` method must be called in a loop
- The `notify` method is called to tell one waiting thread that the event it was waiting has happened
- The `notifyAll` method awakens all of the threads on the object's wait list

Java's Thread Evaluation

- Java's support for concurrency is relatively simple but effective
- Not as powerful as Ada's tasks

C# Threads

- Loosely based on Java but there are significant differences
- Basic thread operations
 - ▣ Any method can run in its own thread
 - ▣ A thread is created by creating a `Thread` object
 - ▣ Creating a thread does not start its concurrent execution; it must be requested through the `Start` method
 - ▣ A thread can be made to wait for another thread to finish with `Join`
 - ▣ A thread can be suspended with `Sleep`
 - ▣ A thread can be terminated with `Abort`

Synchronizing Threads

- Three ways to synchronize C# threads
 - ▣ The `Interlocked` class
 - Used when the only operations that need to be synchronized are incrementing or decrementing of an integer
 - ▣ The `lock` statement
 - Used to mark a critical section of code in a thread
`lock (expression) { ... }`
 - ▣ The `Monitor` class
 - Provides four methods that can be used to provide more sophisticated synchronization

C#'s Concurrency Evaluation

- An advance over Java threads, e.g., any method can run its own thread
- Thread termination is cleaner than in Java
- Synchronization is more sophisticated

Statement-Level Concurrency

- Objective: Provide a mechanism that the programmer can use to inform compiler of ways it can map the program onto multiprocessor architecture
- Minimize communication among processors and the memories of the other processors

High-Performance Fortran

- A collection of extensions that allow the programmer to provide information to the compiler to help it optimize code for multiprocessor computers
- Specify the number of processors, the distribution of data over the memories of those processors, and the alignment of data

Primary HPF Specifications

- Number of processors

`!HPF$ PROCESSORS procs (n)`

- Distribution of data

`!HPF$ DISTRIBUTE (kind) ONTO procs ::
 identifier_list`

- ▣ *kind* can be BLOCK (distribute data to processors in blocks) or CYCLIC (distribute data to processors one element at a time)

- Relate the distribution of one array with that of another

`ALIGN array1_element WITH array2_element`

Statement-Level Concurrency Example

```
REAL list_1(1000), list_2(1000)
INTEGER list_3(500), list_4(501)
!HPF$ PROCESSORS proc (10)
!HPF$ DISTRIBUTE (BLOCK) ONTO procs ::
                                list_1, list_2
!HPF$ ALIGN list_1(index) WITH
                                list_4 (index+1)
...
list_1 (index) = list_2(index)
list_3(index) = list_4(index+1)
```

Statement-Level Concurrency (continued)

- **FORALL** statement is used to specify a list of statements that may be executed concurrently

```
FORALL (index = 1:1000)
```

```
    list_1(index) = list_2(index)
```

- Specifies that all 1,000 RHSs of the assignments can be evaluated before any assignment takes place